

AGENTS.md

This repository uses AI engineering agents (Codex, Antigravity and others) under the supervision of a principal engineering GPT.

Agents must follow strict engineering practices.

Core Principles

1. Clarity over cleverness
2. Maintainable code over fast hacks
3. Security by default
4. Automation where possible
5. Observability in production systems
6. Minimal architectural complexity

Responsibilities of Agents

Agents may:

- implement features
- refactor modules
- write tests
- improve documentation
- debug errors
- propose optimizations

Agents must NOT:

- change architecture without explicit instruction
- introduce new dependencies without justification
- modify unrelated modules
- remove security validations
- bypass tests

Code Quality Rules

All generated code must:

- be readable
- use clear naming
- avoid duplication
- respect separation of concerns
- include error handling
- include validation where necessary
- follow repository coding standards

When a task is ambiguous

The agent must:

1. state assumptions
2. propose alternatives
3. ask for clarification if needed

Never guess silently.

Testing Requirements

Every non-trivial change must include:

- unit tests
- validation of existing functionality
- regression prevention

Security Requirements

Agents must consider:

- authentication
- authorization
- input validation
- secret management
- tenant isolation when applicable

Security shortcuts are forbidden.

Engineering Philosophy

Prefer simple robust solutions that evolve well over time.

Avoid premature optimization.

Avoid unnecessary complexity.